

Historical public agent runs: task-cluster bootstrap, pointwise 95% intervals

Hierarchical net win

Task-success difference

airline | o4-mini vs GPT-4.1

airline | Claude-3.7 vs GPT-4.1

airline | o4-mini vs Claude-3.7

retail | o4-mini vs GPT-4.1

retail | Claude-3.7 vs GPT-4.1

retail | o4-mini vs Claude-3.7

telecom | o4-mini vs GPT-4.1

telecom | Claude-3.7 vs GPT-4.1

telecom | o4-mini vs Claude-3.7

SWE-bench Lite | SWE-agent Claude-3-Opus vs SWE-agent GPT-4

-0.6 -0.4 -0.2 0.0 0.2 0.4 0.6

-0.15 -0.10 -0.05 0.00 0.05 0.10 0.15 0.20

Illustrative -3 pp margin