


Guarded betting e-process, o4-mini vs GPT-4.1 (tau2-bench, all domains): time to deployment

Replicates deployed by n



Paired (shadow) design
95.2% deployed

Cross-arrival design
36.0% deployed

Canonical run (replay_results.csv, 500 orders):

paired: 94.2% deployed, median stop 852 (IQR 588–1148), harm flags 0.0%

cross-arrival: 35.4% deployed, median stop 1204 (IQR 734–1573), harm flags 0.2%

Curves: independent re-run of the same replay (fresh seeds); dotted = re-run median stop