


Guarded betting e-process, o4-mini vs GPT-4.1 (tau2-bench, all domains): time to deployment

Replicates deployed by n



Canonical run (replay_results.csv, 500 orders):

paired: 97.6% deployed, median stop 838 (IQR 567–1198), harm flags 0.0%

cross-arrival: 36.8% deployed, median stop 1228 (IQR 861–1646), harm flags 0.2%

Curves: independent re-run of the same replay (fresh seeds); dotted = re-run median stop